

DATA606_Lab5a_RJM

RJM

2019-12-29

In this lab, we investigate the ways in which the statistics from a random sample of data can serve as point estimates for population parameters. We're interested in formulating a *sampling distribution* of our estimate in order to learn about the properties of the estimate, such as its distribution.

The data

We consider real estate data from the city of Ames, Iowa. The details of every real estate transaction in Ames is recorded by the City Assessor's office. Our particular focus for this lab will be all residential home sales in Ames between 2006 and 2010. This collection represents our population of interest. In this lab we would like to learn about these home sales by taking smaller samples from the full population. Let's load the data.

```
download.file("http://www.openintro.org/stat/data/ames.RData", destfile = "ames.RData")
load("ames.RData")
head(ames)
```

##	Order	PID	MS.SubClass	MS.Zoning	Lot.Frontage	Lot.Area	Street	Alley
## 1	1	526301100	20	RL	141	31770	Pave	<NA>
## 2	2	526350040	20	RH	80	11622	Pave	<NA>
## 3	3	526351010	20	RL	81	14267	Pave	<NA>
## 4	4	526353030	20	RL	93	11160	Pave	<NA>
## 5	5	527105010	60	RL	74	13830	Pave	<NA>
## 6	6	527105030	60	RL	78	9978	Pave	<NA>
##	Lot.Shape	Land.Contour	Utilities	Lot.Config	Land.Slope	Neighborhood		
## 1	IR1	Lvl	AllPub	Corner	Gtl	Names		
## 2	Reg	Lvl	AllPub	Inside	Gtl	Names		
## 3	IR1	Lvl	AllPub	Corner	Gtl	Names		
## 4	Reg	Lvl	AllPub	Corner	Gtl	Names		
## 5	IR1	Lvl	AllPub	Inside	Gtl	Gilbert		
## 6	IR1	Lvl	AllPub	Inside	Gtl	Gilbert		
##	Condition.1	Condition.2	Bldg.Type	House.Style	Overall.Qual	Overall.Cond		
## 1	Norm	Norm	1Fam	1Story	6	5		
## 2	Feedr	Norm	1Fam	1Story	5	6		
## 3	Norm	Norm	1Fam	1Story	6	6		
## 4	Norm	Norm	1Fam	1Story	7	5		
## 5	Norm	Norm	1Fam	2Story	5	5		
## 6	Norm	Norm	1Fam	2Story	6	6		
##	Year.Built	Year.Remod.Add	Roof.Style	Roof.Matl	Exterior.1st	Exterior.2nd		
## 1	1960	1960	Hip	CompShg	BrkFace	Plywood		
## 2	1961	1961	Gable	CompShg	VinylSd	VinylSd		

## 3	1958	1958	Hip	CompShg	Wd Sdng	Wd Sdng
## 4	1968	1968	Hip	CompShg	BrkFace	BrkFace
## 5	1997	1998	Gable	CompShg	VinylSd	VinylSd
## 6	1998	1998	Gable	CompShg	VinylSd	VinylSd
##	Mas.Vnr.Type	Mas.Vnr.Area	Exter.Qual	Exter.Cond	Foundation	Bsmt.Qual
## 1	Stone	112	TA	TA	CBlock	TA
## 2	None	0	TA	TA	CBlock	TA
## 3	BrkFace	108	TA	TA	CBlock	TA
## 4	None	0	Gd	TA	CBlock	TA
## 5	None	0	TA	TA	PConc	Gd
## 6	BrkFace	20	TA	TA	PConc	TA
##	Bsmt.Cond	Bsmt.Exposure	BsmtFin.Type.1	BsmtFin.SF.1	BsmtFin.Type.2	
## 1	Gd	Gd	BLQ	639	Unf	
## 2	TA	No	Rec	468	LwQ	
## 3	TA	No	ALQ	923	Unf	
## 4	TA	No	ALQ	1065	Unf	
## 5	TA	No	GLQ	791	Unf	
## 6	TA	No	GLQ	602	Unf	
##	BsmtFin.SF.2	Bsmt.Unf.SF	Total.Bsmt.SF	Heating	Heating.QC	Central.Air
## 1	0	441	1080	GasA	Fa	Y
## 2	144	270	882	GasA	TA	Y
## 3	0	406	1329	GasA	TA	Y
## 4	0	1045	2110	GasA	Ex	Y
## 5	0	137	928	GasA	Gd	Y
## 6	0	324	926	GasA	Ex	Y
##	Electrical	X1st.Flr.SF	X2nd.Flr.SF	Low.Qual.Fin.SF	Gr.Liv.Area	Bsmt.Full.Bath
## 1	SBrkr	1656	0	0	1656	1
## 2	SBrkr	896	0	0	896	0
## 3	SBrkr	1329	0	0	1329	0
## 4	SBrkr	2110	0	0	2110	1
## 5	SBrkr	928	701	0	1629	0
## 6	SBrkr	926	678	0	1604	0
##	Bsmt.Half.Bath	Full.Bath	Half.Bath	Bedroom.AbvGr	Kitchen.AbvGr	Kitchen.Qual
## 1	0	1	0	3	1	TA
## 2	0	1	0	2	1	TA
## 3	0	1	1	3	1	Gd
## 4	0	2	1	3	1	Ex
## 5	0	2	1	3	1	TA
## 6	0	2	1	3	1	Gd
##	TotRms.AbvGrd	Functional	Fireplaces	Fireplace.Qu	Garage.Type	Garage.Yr.Blt
## 1	7	Typ	2	Gd	Attchd	1960
## 2	5	Typ	0	<NA>	Attchd	1961
## 3	6	Typ	0	<NA>	Attchd	1958
## 4	8	Typ	2	TA	Attchd	1968
## 5	6	Typ	1	TA	Attchd	1997
## 6	7	Typ	1	Gd	Attchd	1998
##	Garage.Finish	Garage.Cars	Garage.Area	Garage.Qual	Garage.Cond	Paved.Drive
## 1	Fin	2	528	TA	TA	P
## 2	Unf	1	730	TA	TA	Y
## 3	Unf	1	312	TA	TA	Y
## 4	Fin	2	522	TA	TA	Y
## 5	Fin	2	482	TA	TA	Y
## 6	Fin	2	470	TA	TA	Y
##	Wood.Deck.SF	Open.Porch.SF	Enclosed.Porch	X3Ssn.Porch	Screen.Porch	Pool.Area

```
## 1      210      62      0      0      0      0
## 2      140       0      0      0      120     0
## 3      393      36      0      0      0      0
## 4       0       0      0      0      0      0
## 5      212      34      0      0      0      0
## 6      360      36      0      0      0      0
##   Pool.QC Fence Misc.Feature Misc.Val Mo.Sold Yr.Sold Sale.Type Sale.Condition
## 1   <NA> <NA>      <NA>      0      5    2010      WD      Normal
## 2   <NA> MnPrv      <NA>      0      6    2010      WD      Normal
## 3   <NA> <NA>      Gar2    12500      6    2010      WD      Normal
## 4   <NA> <NA>      <NA>      0      4    2010      WD      Normal
## 5   <NA> MnPrv      <NA>      0      3    2010      WD      Normal
## 6   <NA> <NA>      <NA>      0      6    2010      WD      Normal
##   SalePrice
## 1    215000
## 2    105000
## 3    172000
## 4    244000
## 5    189900
## 6    195500
```

We see that there are quite a few variables in the data set, enough to do a very in-depth analysis. For this lab, we'll restrict our attention to just two of the variables: the above ground living area of the house in square feet (**Gr.Liv.Area**) and the sale price (**SalePrice**). To save some effort throughout the lab, create two variables with short names that represent these two variables.

```
area <- ames$Gr.Liv.Area
price <- ames$SalePrice
```


Let's look at the distribution of area in our population of home sales by calculating a few summary statistics and making a histogram.

```
summary(area)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      334   1126   1442   1500   1743   5642
```

```
hist(area)
```

Histogram of area



```
boxplot(area)
```



1. Describe this population distribution. The distribution is unimodal with a very slight right skew. ##
The unknown sampling distribution

In this lab we have access to the entire population, but this is rarely the case in real life. Gathering information on an entire population is often extremely costly or impossible. Because of this, we often take a sample of the population and use that to understand the properties of the population.


If we were interested in estimating the mean living area in Ames based on a sample, we can use the following command to survey the population.

```
set.seed(123)
samp1 <- sample(area, 50)
summary(samp1)
```


```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      767   1116   1398   1459   1680   2610
```

```
hist(samp1)
```

Histogram of samp1



```
boxplot(samp1)
```



This command collects a simple random sample of size 50 from the vector `area`, which is assigned to `samp1`. This is like going into the City Assessor's database and pulling up the files on 50 random home sales. Working with these 50 files would be considerably simpler than working with all 2930 home sales.

2. Describe the distribution of this sample. How does it compare to the distribution of the population? The distribution is still unimodal but the skew is a bit less probably a result of a smaller set. There are also some possible outliers.

If we're interested in estimating the average living area in homes in Ames using the sample, our best single guess is the sample mean.

```
mean(samp1)
```

```
## [1] 1459.08
```

Depending on which 50 homes you selected, your estimate could be a bit above or a bit below the true population mean of 1499.69 square feet. In general, though, the sample mean turns out to be a pretty good estimate of the average living area, and we were able to get it by sampling less than 3% of the population.

3. Take a second sample, also of size 50, and call it `samp2`. How does the mean of `samp2` compare with the mean of `samp1`? Suppose we took two more samples, one of size 100 and one of size 1000. Which would you think would provide a more accurate estimate of the population mean? If the seed is not set or set at a different point, the mean is slightly varied for both samples but lies closely to the population mean in both cases. If the seed is different or not set then the mean is slightly varied. The mean moved a bit further from the population mean for a random sample of 100 but then fell closer to the population mean at size 1000.

```
set.seed(123)
samp2 <- sample(area, 50)
samp3 <- sample(area, 100)
samp4 <- sample(area, 1000)
```

```
mean(samp2)
```

```
## [1] 1459.08
```

```
mean(samp3)
```

```
## [1] 1469.52
```

```
mean(samp4)
```

```
## [1] 1494.928
```


Not surprisingly, every time we take another random sample, we get a different sample mean. It's useful to get a sense of just how much variability we should expect when estimating the population mean this way. The distribution of sample means, called the *sampling distribution*, can help us understand this variability. In this lab, because we have access to the population, we can build up the sampling distribution for the sample mean by repeating the above steps many times. Here we will generate 5000 samples and compute the sample mean of each.

```
sample_means50 <- rep(NA, 5000)

for(i in 1:5000){
  samp <- sample(area, 50)
  sample_means50[i] <- mean(samp)
}

hist(sample_means50)
```


Histogram of sample_means50



If you would like to adjust the bin width of your histogram to show a little more detail, you can do so by changing the `breaks` argument.

```
hist(sample_means50, breaks = 25)
```

Histogram of sample_means50



Here we use R to take 5000 samples of size 50 from the population, calculate the mean of each sample, and store each result in a vector called `sample_means50`. On the next page, we'll review how this set of code works.

4. How many elements are there in `sample_means50`? Describe the sampling distribution, and be sure to specifically note its center. Would you expect the distribution to change if we instead collected 50,000 sample means? There are 5,000 elements in `sample_means50` with each element having 50 samples. The distribution looks normal and the mean is very close to 1,500 which in turn is very close to the population mean of 1,499.69. The distribution will still be normal if we increase the number of elements to 50,000 as we are increasing the number of our samples. Also, the mean should still be close to the population mean of 1,499.69.

Interlude: The for loop

Let's take a break from the statistics for a moment to let that last block of code sink in. You have just run your first `for` loop, a cornerstone of computer programming. The idea behind the `for` loop is *iteration*: it allows you to execute code as many times as you want without having to type out every iteration. In the case above, we wanted to iterate the two lines of code inside the curly braces that take a random sample of size 50 from `area` then save the mean of that sample into the `sample_means50` vector. Without the `for` loop, this would be painful:

```
sample_means50 <- rep(NA, 5000)

samp <- sample(area, 50)
sample_means50[1] <- mean(samp)

samp <- sample(area, 50)
sample_means50[2] <- mean(samp)
```



```
samp <- sample(area, 50)
sample_means50[3] <- mean(samp)

samp <- sample(area, 50)
sample_means50[4] <- mean(samp)
```

and so on...

With the for loop, these thousands of lines of code are compressed into a handful of lines. We've added one extra line to the code below, which prints the variable `i` during each iteration of the for loop. Run this code.

```
sample_means50 <- rep(NA, 5000)

for(i in 1:5000){
  samp <- sample(area, 50)
  sample_means50[i] <- mean(samp)
  # print(i)
}
```

Let's consider this code line by line to figure out what it does. In the first line we *initialized a vector*. In this case, we created a vector of 5000 zeros called `sample_means50`. This vector will store values generated within the for loop.

The second line calls the for loop itself. The syntax can be loosely read as, "for every element `i` from 1 to 5000, run the following lines of code". You can think of `i` as the counter that keeps track of which loop you're on. Therefore, more precisely, the loop will run once when `i = 1`, then once when `i = 2`, and so on up to `i = 5000`.

The body of the for loop is the part inside the curly braces, and this set of code is run for each value of `i`. Here, on every loop, we take a random sample of size 50 from `area`, take its mean, and store it as the `i`th element of `sample_means50`.

In order to display that this is really happening, we asked R to print `i` at each iteration. This line of code is optional and is only used for displaying what's going on while the for loop is running.

The for loop allows us to not just run the code 5000 times, but to neatly package the results, element by element, into the empty vector that we initialized at the outset.

5. To make sure you understand what you've done in this loop, try running a smaller version. Initialize a vector of 100 zeros called `sample_means_small`. Run a loop that takes a sample of size 50 from `area` and stores the sample mean in `sample_means_small`, but only iterate from 1 to 100. Print the output to your screen (type `sample_means_small` into the console and press enter). How many elements are there in this object called `sample_means_small`? What does each element represent?

```
sample_means_small <- rep(NA, 100)

for(i in 1:100){
  samp <- sample(area, 50)
  sample_means_small[i] <- mean(samp)
}

sample_means_small
```


```
## [1] 1409.88 1455.14 1483.06 1401.56 1416.70 1501.24 1530.90 1590.58 1525.72
## [10] 1400.40 1490.00 1549.32 1520.52 1518.38 1464.92 1415.50 1507.80 1465.44
## [19] 1488.42 1523.06 1474.78 1557.62 1459.46 1511.72 1411.52 1584.58 1487.64
## [28] 1504.72 1511.68 1503.32 1459.88 1349.40 1488.12 1441.66 1462.44 1457.86
## [37] 1461.00 1402.34 1413.34 1485.94 1597.72 1602.00 1646.04 1694.24 1479.00
## [46] 1349.48 1331.84 1363.90 1516.62 1444.06 1499.70 1500.30 1390.10 1539.74
## [55] 1602.30 1515.72 1569.86 1454.04 1622.82 1521.34 1424.20 1491.64 1531.38
## [64] 1541.34 1474.26 1489.10 1658.76 1559.80 1408.94 1526.14 1461.60 1517.00
## [73] 1525.72 1451.12 1635.00 1578.94 1465.66 1377.46 1531.16 1486.30 1514.30
## [82] 1482.28 1489.24 1487.30 1474.30 1383.26 1596.02 1484.46 1529.82 1656.74
## [91] 1357.90 1453.54 1451.78 1573.42 1475.52 1573.68 1393.72 1531.78 1367.96
## [100] 1559.98
```

There are 100 elements in this object

Sample size and the sampling distribution

Mechanics aside, let's return to the reason we used a `for` loop: to compute a sampling distribution, specifically, this one.

```
hist(sample_means50)
```



The sampling distribution that we computed tells us much about estimating the average living area in homes in Ames. Because the sample mean is an unbiased estimator, the sampling distribution is centered at the true average living area of the the population, and the spread of the distribution indicates how much variability is induced by sampling only 50 home sales.

To get a sense of the effect that sample size has on our distribution, let's build up two more sampling distributions: one based on a sample size of 10 and another based on a sample size of 100.

```

sample_means10 <- rep(NA, 5000)
sample_means100 <- rep(NA, 5000)

for(i in 1:5000){
  samp <- sample(area, 10)
  sample_means10[i] <- mean(samp)
  samp <- sample(area, 100)
  sample_means100[i] <- mean(samp)
}

```

Here we're able to use a single `for` loop to build two distributions by adding additional lines inside the curly braces. Don't worry about the fact that `samp` is used for the name of two different objects. In the second command of the `for` loop, the mean of `samp` is saved to the relevant place in the vector `sample_means10`. With the mean saved, we're now free to overwrite the object `samp` with a new sample, this time of size 100. In general, anytime you create an object using a name that is already in use, the old object will get replaced with the new one.

To see the effect that different sample sizes have on the sampling distribution, plot the three distributions on top of one another.

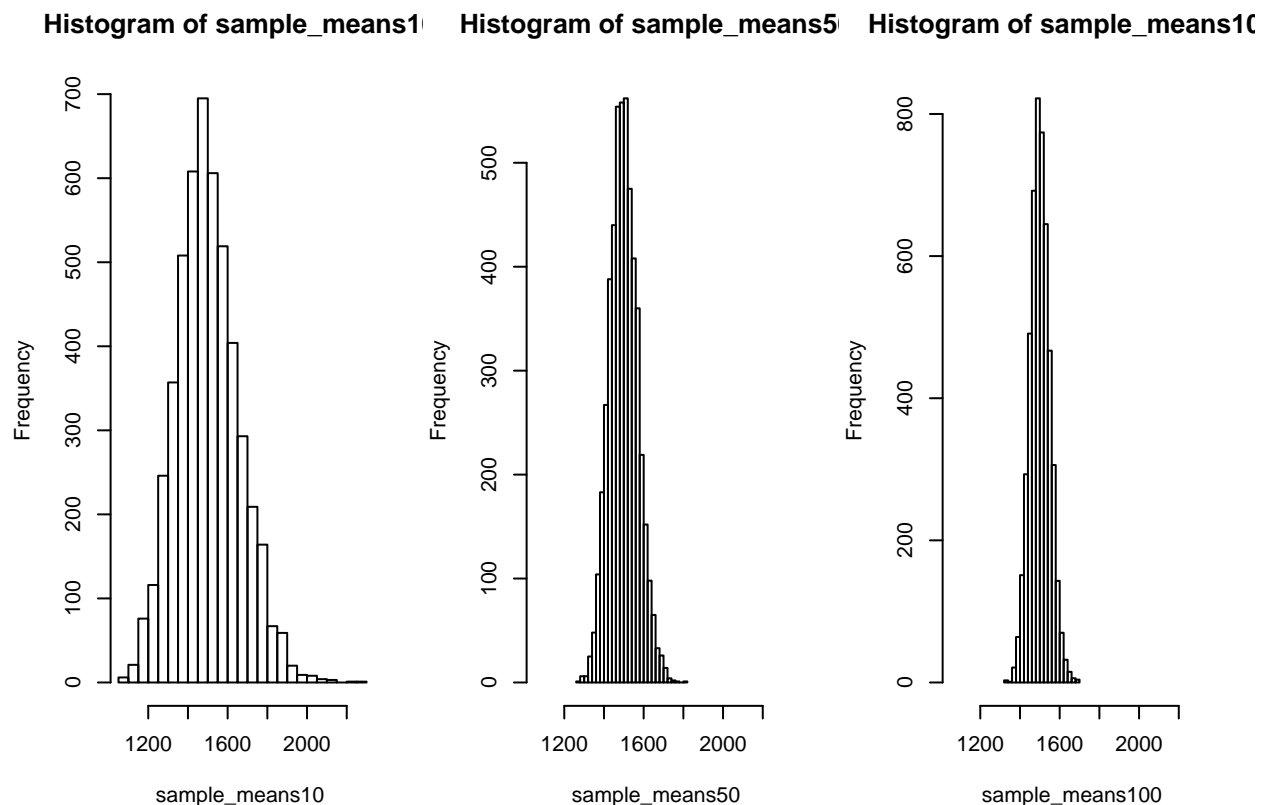
```

par(mfrow = c(1, 3))

xlimits <- range(sample_means10)

hist(sample_means10, breaks = 20, xlim = xlimits)
hist(sample_means50, breaks = 20, xlim = xlimits)
hist(sample_means100, breaks = 20, xlim = xlimits)

```



The first command specifies that you'd like to divide the plotting area into 3 rows and 1 column of plots (to return to the default setting of plotting one at a time, use `par(mfrow = c(1, 1))`). The `breaks` argument specifies the number of bins used in constructing the histogram. The `xlim` argument specifies the range of the x-axis of the histogram, and by setting it equal to `xlimits` for each histogram, we ensure that all three histograms will be plotted with the same limits on the x-axis.

6. When the sample size is larger, what happens to the center? What about the spread?

```
xlimits_10 <- range(sample_means10)
xlimits_50 <- range(sample_means50)
xlimits_100 <- range(sample_means100)

xlimits_10[2] - xlimits_10[1]
```

```
## [1] 1207.7
```

```
xlimits_50[2] - xlimits_50[1]
```

```
## [1] 544.16
```

```
xlimits_100[2] - xlimits_100[1]
```

```
## [1] 367.97
```

From the above, the centre is getting closer to the population with an increasing sample size making the spread smaller.

On your own

So far, we have only focused on estimating the mean living area in homes in Ames. Now you'll try to estimate the mean home price.

- Take a random sample of size 50 from `price`. Using this sample, what is your best point estimate of the population mean?

```
set.seed(123)
samp_p <- sample(price, 50)
samp_p_m <- mean(samp_p)
samp_p_m
```

```
## [1] 175685.2
```

```
# The best estimate of population mean is $175,032.10 based on a random sample
# of 50.
```


- Since you have access to the population, simulate the sampling distribution for \bar{x}_{price} by taking 5000 samples from the population of size 50 and computing 5000 sample means. Store these means in a vector called `sample_means50`. Plot the data, then describe the shape of this sampling distribution. Based on this sampling distribution, what would you guess the mean home price of the population to be? Finally, calculate and report the population mean.

```
sample_means50 <- rep(NA, 5000)

for(i in 1:5000){
  samp <- sample(price, 50)
  sample_means50[i] <- mean(samp)
}

hist(sample_means50, breaks = 25)
```

Histogram of sample_means50



From the above histogram, we can see that the sampling distribution is normally distributed with a mean

```
pop_m <- mean(sample_means50)
pop_m
```

```
## [1] 180872.2
```

The population mean is very close to the mean of the 5000 sample means of 5000 random samples of the population of size 50.

- Change your sample size from 50 to 150, then compute the sampling distribution using the same method as above, and store these means in a new vector called `sample_means150`. Describe the shape of this sampling distribution, and compare it to the sampling distribution for a sample size of 50. Based on this sampling distribution, what would you guess to be the mean sale price of homes in Ames?

```


sample_means150 <- rep(NA, 5000)

for(i in 1:5000){
  samp <- sample(price, 150)
  sample_means150[i] <- mean(samp)
}

hist(sample_means150, breaks = 25)

```

Histogram of sample_means150



```


par(mfrow = c(1, 2))

xlimits <- range(sample_means50)


hist(sample_means50, breaks = 20, xlim = xlimits)
hist(sample_means150, breaks = 20, xlim = xlimits)

```

Histogram of sample_means50



Histogram of sample_means150



From the above, we can observe that the means have converged around \$180,000 and the spread has become smaller as the sample size has been increased from 50 to 150.

- Of the sampling distributions from 2 and 3, which has a smaller spread? If we're concerned with making estimates that are more often close to the true value, would we prefer a distribution with a large or small spread?

```
xlimits50 <- range(sample_means50)
xlimits150 <- range(sample_means150)

xlimits50[2] - xlimits50[1]
```

```
## [1] 77613.26
```

```
xlimits150[2] - xlimits150[1]
```

```
## [1] 49079.03
```

A distribution with smaller spread is preferred as there is less standard deviation from the mean which means that more of the observed values are closer to the mean.