

DATA606_Lab8_RJM

RJM

2019-12-29

Batter up

The movie Moneyball focuses on the “quest for the secret of success in baseball”. It follows a low-budget team, the Oakland Athletics, who believed that underused statistics, such as a player’s ability to get on base, better predict the ability to score runs than typical statistics like home runs, RBIs (runs batted in), and batting average. Obtaining players who excelled in these underused statistics turned out to be much more affordable for the team.

In this lab we’ll be looking at data from all 30 Major League Baseball teams and examining the linear relationship between runs scored in a season and a number of other player statistics. Our aim will be to summarize these relationships both graphically and numerically in order to find which variable, if any, helps us best predict a team’s runs scored in a season.

The data

Let’s load up the data for the 2011 season.

```
download.file("http://www.openintro.org/stat/data/mlb11.RData", destfile = "mlb11.RData")

load("mlb11.RData")


head(mlb11)
```

##		team	runs	at_bats	hits	homeruns	bat_avg	strikeouts
## 1		Texas Rangers	855	5659	1599	210	0.283	930
## 2		Boston Red Sox	875	5710	1600	203	0.280	1108
## 3		Detroit Tigers	787	5563	1540	169	0.277	1143
## 4		Kansas City Royals	730	5672	1560	129	0.275	1006
## 5		St. Louis Cardinals	762	5532	1513	162	0.273	978
## 6		New York Mets	718	5600	1477	108	0.264	1085
##		stolen_bases	wins	new_onbase	new_slug	new_obs		
## 1		143	96	0.340	0.460	0.800		
## 2		102	90	0.349	0.461	0.810		
## 3		49	95	0.340	0.434	0.773		
## 4		153	71	0.329	0.415	0.744		
## 5		57	90	0.341	0.425	0.766		
## 6		130	77	0.335	0.391	0.725		

In addition to runs scored, there are seven traditionally used variables in the data set: at-bats, hits, home runs, batting average, strikeouts, stolen bases, and wins. There are also three newer variables: on-base percentage, slugging percentage, and on-base plus slugging. For the first portion of the analysis we’ll consider the seven traditional variables. At the end of the lab, you’ll work with the newer variables on your own.

1. What type of plot would you use to display the relationship between `runs` and one of the other numerical variables? Plot this relationship using the variable `at_bats` as the predictor. Does the relationship look linear? If you knew a team's `at_bats`, would you be comfortable using a linear model to predict the number of runs?

```
plot(mlb11$at_bats, mlb11$runs)
abline(lm(mlb11$runs ~ mlb11$at_bats, data = mlb11), col = "blue")
```



The relationship looks linear (not strong) as most points on graph are converged around the classifier line with very few outliers.

If the relationship looks linear, we can quantify the strength of the relationship with the correlation coefficient.

```
cor(mlb11$runs, mlb11$at_bats)
```

```
## [1] 0.610627
```

Sum of squared residuals


Think back to the way that we described the distribution of a single variable. Recall that we discussed characteristics such as center, spread, and shape. It's also useful to be able to describe the relationship of two numerical variables, such as `runs` and `at_bats` above.

2. Looking at your plot from the previous exercise, describe the relationship between these two variables. Make sure to discuss the form, direction, and strength of the relationship as well as any unusual observations.

The relationship appears to be linear but not a strong one. The correlation coefficient is 0.61 and most of the points appear to conform to the linearity but there are a few outliers.

Just as we used the mean and standard deviation to summarize a single variable, we can summarize the relationship between these two variables by finding the line that best follows their association. Use the following interactive function to select the line that you think does the best job of going through the cloud of points.

```
plot_ss(x = mlb11$at_bats, y = mlb11$runs)
```




```
## Click two points to make a line.
## Call:
## lm(formula = y ~ x, data = pts)
##
## Coefficients:
## (Intercept)          x
## -2789.2429      0.6305
##
## Sum of Squares: 123721.9
```

After running this command, you'll be prompted to click two points on the plot to define a line. Once you've done that, the line you specified will be shown in black and the residuals in blue. Note that there are 30 residuals, one for each of the 30 observations. Recall that the residuals are the difference between the observed values and the values predicted by the line:

$$e_i = y_i - \hat{y}_i$$

The most common way to do linear regression is to select the line that minimizes the sum of squared residuals. To visualize the squared residuals, you can rerun the plot command and add the argument `showSquares = TRUE`.

```
plot_ss(x = mlb11$at_bats, y = mlb11$runs, showSquares = TRUE)
```



```
## Click two points to make a line.
## Call:
## lm(formula = y ~ x, data = pts)
##
## Coefficients:
## (Intercept)          x
## -2789.2429      0.6305
##
## Sum of Squares: 123721.9
```


Note that the output from the `plot_ss` function provides you with the slope and intercept of your line as well as the sum of squares.

- Using `plot_ss`, choose a line that does a good job of minimizing the sum of squares. Run the function several times. What was the smallest sum of squares that you got? How does it compare to your neighbors?

```
plot_ss(x = mlb11$at_bats, y = mlb11$runs, showSquares = TRUE)
```

```
## Click two points to make a line.
## Call:
## lm(formula = y ~ x, data = pts)
##
## Coefficients:
## (Intercept)          x
## -2789.2429      0.6305
##
## Sum of Squares: 123721.9
```

```
plot_ss(x = mlb11$at_bats, y = mlb11$runs, showSquares = TRUE)
```



```
## Click two points to make a line.
## Call:
## lm(formula = y ~ x, data = pts)
##
## Coefficients:
## (Intercept)          x
## -2789.2429      0.6305
##
## Sum of Squares:  123721.9
```

```
plot_ss(x = mlb11$at_bats, y = mlb11$runs, showSquares = TRUE)
```

```
## Click two points to make a line.
## Call:
## lm(formula = y ~ x, data = pts)
##
## Coefficients:
## (Intercept)          x
## -2789.2429      0.6305
##
## Sum of Squares:  123721.9
```

The least sum of squares appears to be 123721.90 which is same for all the lines.

The linear model

It is rather cumbersome to try to get the correct least squares line, i.e. the line that minimizes the sum of squared residuals, through trial and error. Instead we can use the `lm` function in R to fit the linear model (a.k.a. regression line).

```
m1 <- lm(runs ~ at_bats, data = mlb11)
```

The first argument in the function `lm` is a formula that takes the form $y \sim x$. Here it can be read that we want to make a linear model of `runs` as a function of `at_bats`. The second argument specifies that R should look in the `mlb11` data frame to find the `runs` and `at_bats` variables.

The output of `lm` is an object that contains all of the information we need about the linear model that was just fit. We can access this information using the summary function.

```
summary(m1)
```

```
##
## Call:
## lm(formula = runs ~ at_bats, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -125.58  -47.05  -16.59   54.40  176.87
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -2789.2429   853.6957  -3.267 0.002871 **
## at_bats       0.6305     0.1545   4.080 0.000339 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 66.47 on 28 degrees of freedom
## Multiple R-squared:  0.3729, Adjusted R-squared:  0.3505
## F-statistic: 16.65 on 1 and 28 DF,  p-value: 0.0003388
```

Let's consider this output piece by piece. First, the formula used to describe the model is shown at the top. After the formula you find the five-number summary of the residuals. The “Coefficients” table shown next is key; its first column displays the linear model's y-intercept and the coefficient of `at_bats`. With this table, we can write down the least squares regression line for the linear model:

$$\hat{y} = -2789.2429 + 0.6305 * atbats$$

One last piece of information we will discuss from the summary output is the Multiple R-squared, or more simply, R^2 . The R^2 value represents the proportion of variability in the response variable that is explained by the explanatory variable. For this model, 37.3% of the variability in runs is explained by at-bats.

4. Fit a new model that uses `homeruns` to predict `runs`. Using the estimates from the R output, write the equation of the regression line. What does the slope tell us in the context of the relationship between success of a team and its home runs?

```
m2 <- lm(runs ~ homeruns, data = mlb11)
```

```
summary(m2)
```

```
##
## Call:
```


```
## lm(formula = runs ~ homeruns, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -91.615 -33.410   3.231  24.292 104.631
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  415.2389    41.6779   9.963 1.04e-10 ***
## homeruns      1.8345     0.2677   6.854 1.90e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 51.29 on 28 degrees of freedom
## Multiple R-squared:  0.6266, Adjusted R-squared:  0.6132
## F-statistic: 46.98 on 1 and 28 DF,  p-value: 1.9e-07
```

$y_{\text{hat}} = 415.24 + 1.83 * \text{homeruns}$ From the above, we can see that runs go up by 1.83 with an increase of 1 in homeruns. Also, 62.66% of the variability in runs is explained by homeruns which indicates a strong correlation.

Prediction and prediction errors

Let's create a scatterplot with the least squares line laid on top.

```
plot(mlb11$runs ~ mlb11$at_bats)
abline(m1)
```



The function `abline` plots a line based on its slope and intercept. Here, we used a shortcut by providing the model `m1`, which contains both parameter estimates. This line can be used to predict y at any value of x . When predictions are made for values of x that are beyond the range of the observed data, it is referred to

as *extrapolation* and is not usually recommended. However, predictions made within the range of the data are more reliable. They're also used to compute the residuals.

5. If a team manager saw the least squares regression line and not the actual data, how many runs would he or she predict for a team with 5,578 at-bats? Is this an overestimate or an underestimate, and by how much? In other words, what is the residual for this prediction?

```
runs <- -2789.2429 + 0.6305 * 5578
paste('The runs with 5578 at_bats are', round(runs, 2))
```


```
## [1] "The runs with 5578 at_bats are 727.69"
```

Since this is the value on the regression line, the residual would be zero. `##` Model diagnostics

To assess whether the linear model is reliable, we need to check for (1) linearity, (2) nearly normal residuals, and (3) constant variability.

Linearity: You already checked if the relationship between runs and at-bats is linear using a scatterplot. We should also verify this condition with a plot of the residuals vs. at-bats. Recall that any code following a `#` is intended to be a comment that helps understand the code but is ignored by R.

```
plot(m1$residuals ~ mlb11$at_bats)
abline(h = 0, lty = 3) # adds a horizontal dashed line at y = 0
```




6. Is there any apparent pattern in the residuals plot? What does this indicate about the linearity of the relationship between runs and at-bats?

There appears to be no pattern, so the linearity does not seem to be positive.


Nearly normal residuals: To check this condition, we can look at a histogram


```
hist(m1$residuals)
```



or a normal probability plot of the residuals.

```
qqnorm(m1$residuals)
qqline(m1$residuals) # adds diagonal line to the normal prob plot
```



7. Based on the histogram and the normal probability plot, does the nearly normal residuals condition appear to be met?

Most points in the histogram appear to be under a normal bell curve and the normal probability plot has most of the points close to the line indicating that the conditions for nearly normal residual condition appear to be met.

Constant variability:

8. Based on the plot in (1), does the constant variability condition appear to be met?

Since the points are above and below the zero line, the constant variability condition is met.

On Your Own

- Choose another traditional variable from `mlb11` that you think might be a good predictor of `runs`. Produce a scatterplot of the two variables and fit a linear model. At a glance, does there seem to be a linear relationship?

```
head(mlb11)
```

```
##           team runs at_bats hits homeruns bat_avg strikeouts
## 1   Texas Rangers  855   5659 1599      210   0.283         930
## 2   Boston Red Sox  875   5710 1600      203   0.280        1108
## 3   Detroit Tigers  787   5563 1540      169   0.277        1143
```

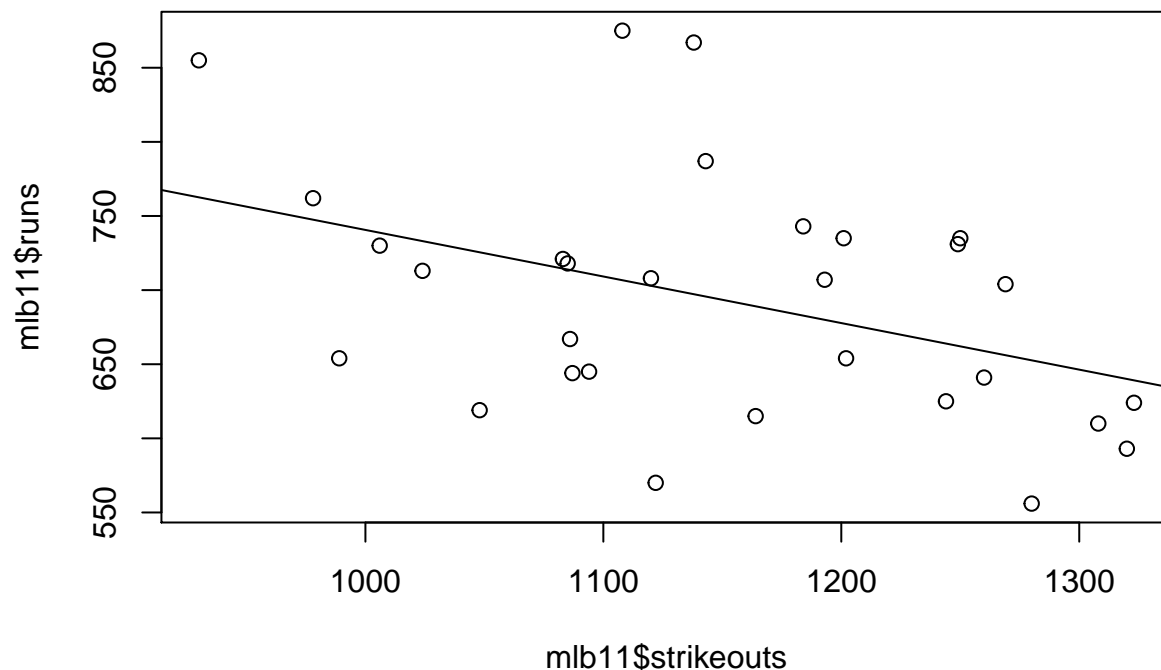
```
## 4 Kansas City Royals 730 5672 1560 129 0.275 1006
## 5 St. Louis Cardinals 762 5532 1513 162 0.273 978
## 6 New York Mets 718 5600 1477 108 0.264 1085
## stolen_bases wins new_onbase new_slug new_obs
## 1 143 96 0.340 0.460 0.800
## 2 102 90 0.349 0.461 0.810
## 3 49 95 0.340 0.434 0.773
## 4 153 71 0.329 0.415 0.744
## 5 57 90 0.341 0.425 0.766
## 6 130 77 0.335 0.391 0.725
```

```
m3 <- lm(runs ~ strikeouts, data = mlb11)
```

```
summary(m3)
```

```
##
## Call:
## lm(formula = runs ~ strikeouts, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -132.27  -46.95  -11.92   55.14  169.76
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1054.7342   151.7890   6.949 1.49e-07 ***
## strikeouts   -0.3141    0.1315  -2.389  0.0239 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 76.5 on 28 degrees of freedom
## Multiple R-squared:  0.1694, Adjusted R-squared:  0.1397
## F-statistic: 5.709 on 1 and 28 DF, p-value: 0.02386
```

```
plot(mlb11$runs ~ mlb11$strikeouts)
abline(m3)
```



Yes, there appears to be a negative linearity between runs and strikeouts but not a strong one.

- How does this relationship compare to the relationship between **runs** and **at_bats**? Use the R^2 values from the two model summaries to compare. Does your variable seem to predict **runs** better than **at_bats**? How can you tell?

The r-squared for variable of strikeouts is 16.94% which tells is that this variable does not explain the variability in runs as better as at-bats which has r-squared valued at 37.29%. From the same logic, at_bats predicts the runs better than the strikeouts.

- Now that you can summarize the linear relationship between two variables, investigate the relationships between **runs** and each of the other five traditional variables. Which variable best predicts **runs**? Support your conclusion using the graphical and numerical methods we've discussed (for the sake of conciseness, only include output for the best variable, not all five).

```
best_match <- cor(mlb11 %>% select(-team, ))[, "runs", drop = F]
best_match
```


```
##               runs
## runs          1.0000000
## at_bats        0.61062705
## hits           0.80121081
## homeruns       0.79155769
## bat_avg        0.80998589
## strikeouts     -0.41153120
## stolen_bases   0.05398141
## wins           0.60080877
## new_onbase     0.92146907
## new_slug       0.94703240
## new_obs        0.96691630
```

```
mlb_bat_avg <- lm(runs ~ bat_avg, data = mlb11)
summary(mlb_bat_avg)
```

```
##
## Call:
## lm(formula = runs ~ bat_avg, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -94.676 -26.303  -5.496   28.482  131.113
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -642.8      183.1   -3.511  0.00153 **
## bat_avg       5242.2      717.3    7.308  5.88e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 49.23 on 28 degrees of freedom
## Multiple R-squared:  0.6561, Adjusted R-squared:  0.6438
## F-statistic: 53.41 on 1 and 28 DF,  p-value: 5.877e-08
```

```
plot(mlb11$runs ~ mlb11$bat_avg)
```

```
abline(mlb_bat_avg)
```



It seems from the above that the variable bat_avg best explains the variance in the runs if we keep the ones starting with “new” aside. The scatter plot and the regression line also seem to agree with this conclusion as most of the point lie close to the line without any clear outlier.

- Now examine the three newer variables. These are the statistics used by the author of *Moneyball* to predict a teams success. In general, are they more or less effective at predicting runs that the old

variables? Explain using appropriate graphical and numerical evidence. Of all ten variables we've analyzed, which seems to be the best predictor of runs? Using the limited (or not so limited) information you know about these baseball statistics, does your result make sense?


```
best_match <- cor(mlb11 %>% select(-team, ))[, "runs", drop = F]
best_match
```

```
##              runs
## runs          1.00000000
## at_bats        0.61062705
## hits           0.80121081
## homeruns       0.79155769
## bat_avg        0.80998589
## strikeouts     -0.41153120
## stolen_bases   0.05398141
## wins           0.60080877
## new_onbase     0.92146907
## new_slug       0.94703240
## new_obs        0.96691630
```

```
mlb_bat_avg2 <- lm(runs ~ new_obs, data = mlb11)
summary(mlb_bat_avg2)
```

```
##
## Call:
## lm(formula = runs ~ new_obs, data = mlb11)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -43.456 -13.690   1.165  13.935  41.156
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -686.61      68.93  -9.962 1.05e-10 ***
## new_obs       1919.36      95.70  20.057 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 21.41 on 28 degrees of freedom
## Multiple R-squared:  0.9349, Adjusted R-squared:  0.9326
## F-statistic: 402.3 on 1 and 28 DF,  p-value: < 2.2e-16
```

```
plot(mlb11$runs ~ mlb11$new_obs)
abline(mlb_bat_avg2)
```

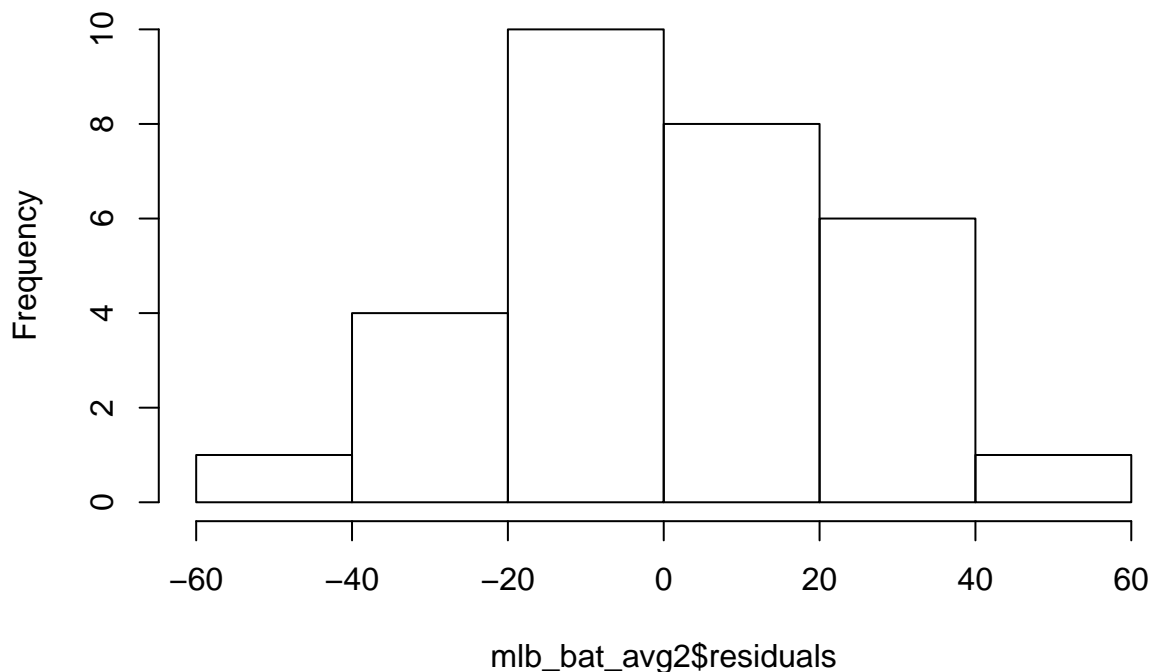


The new_obs variable explains 96.69% of the variance in the runs and the plot seems to agree with the results. Also, the low p-value also points toward a significant explanation of variance through this variable.


- Check the model diagnostics for the regression model with the variable you decided was the best predictor for runs.

```
hist(mlb_bat_avg2$residuals)
```

Histogram of mlb_bat_avg2\$residuals




```
plot(mlb11$runs ~ mlb11$new_obs)  
abline(mlb_bat_avg2)
```



```
qqnorm(mlb_bat_avg2$residuals)  
qqline(mlb_bat_avg2$residuals)
```

Normal Q-Q Plot



The histogram shows a nearly perfect normal distribution. The scatter plot and the regression line also show a very strong linear relationship along with the probability plot that shows a very close fitting for the points with almost no real outliers.